PSI Numerical Methods 2024 - Homework Assignment on Model Fitting & MCMC

We're going to put together everything we have learned so far to re-do the data analysis for the Perlmutter et al. 1999 paper on the discovery of dark energy! (https://ui.adsabs.harvard.edu/abs/1999ApJ...517..565P/abstract (https://ui.adsabs.harvard.edu/abs/1999ApJ...517..565P/abstract))

Start by `Forking` this repository on Github: https://github.com/dstndstn/PSI-Numerical-Methods-2024-MCMC-Homework (https://github.com/dstndstn/PSI-Numerical-Methods-2024-MCMC-Homework) And then clone the repository to your laptop or to Symmetry. You can modify this notebook, and when you are done, save it, and then `git commit -a` the results, and `git push` them back to your fork of the repository. You will "hand in" your homework by giving a link to your Github repository, where the marker will be able to read your notebook.

First, a little bit of background on the cosmology and astrophysics. The paper reports measurements of a group of supernova explosions of a specific type, "Type 1a". These are thought to be caused by a white dwarf star that has a companion star that "donates" gas to the white dwarf. It gradually gains mass until it exceeds the Chandresekhar mass, and explodes. Since they all explode through the same mechanism, and with the same mass, they should all have the same intrinsic brightess. It turns out to be a *little* more complicated than that, but in the end, these Type-1a supernovae can be turned into "standard candles", objects that are all the same brightness. If you can also measure the redshift of each galaxy containing the supernova, then you can map out this brightness--redshift relation, and the shape of that relation depends on how the universe grows over cosmic time. In turn, the growth rate of the universe depends on the contents of the universe!

In this way, these Type-1a supernova allow us to constrain the parameters of a model of the universe. Specifically, the model is called "Lambda-CDM", a universe containing dark energy and matter (cold dark matter, plus regular matter). We will consider a two-parameter version of this model: Ω_M, the amount of matter, and Ω_Λ, the amount of dark energy. These are in cosmology units of "energy density now relative to the critical density", where the critical density is the energy density you need for the universe to be spatially flat (angles of a large triangle sum to 180 degrees). So $\Omega_M = 1$, $\Omega_\Lambda = 0$ would be a flat universe containing all matter, while $\Omega_M = 0.25, \Omega_\Lambda = 0.5$ would be a spatially closed universe with dark energy and matter. Varying these ingredients changes the

growth history of the universe, which changes how much the light from a supernova is redshifted, and how its brightness drops off with distance.

(In the code below, we will call these `Omega_M` = Ω_M and `Omega_DE` = Ω_Λ.)

Distance measurements in cosmology are complicated -- see https://arxiv.org/abs/astro-ph/9905116 (https://arxiv.org/abs/astro-ph/9905116) for details! For this assignment, we will use a cosmology package that will handle all this for us. All we need to use is the "luminosity distance", which is the one that tells you how objects get fainter given a redshift.

```julia
In [1]:
# Let's start by installing the Cosmology package!
using Pkg
Pkg.add("Cosmology")
```

```
        Updating registry at `C:\Users\numbe\.julia\re
   gistries\General.toml`
      Resolving package versions...
    No Changes to `C:\Users\numbe\.julia\environment
   s\v1.10\Project.toml`
      No Changes to `C:\Users\numbe\.julia\environment
   s\v1.10\Manifest.toml`
```

```julia
In [53]:
# We'll also end up using all our old friends:
using WGLMakie
using CSV
using DataFrames
using Cosmology
using Statistics
```

```julia
In [3]:
# There is a data file in this directory, taken basi
data = CSV.read("p99-data.txt", DataFrame, delim=" "
```

```julia
In [4]:
# Make a copy of the data columns that we want to tr
# These are the measured brightnesses, and their Gau
data.mag = data.m_b_eff
data.sigma_mag = data.sigma_m_b_eff;
```

In [5]:
```
f = Figure()
Axis(f[1,1], title="Perlmutter+99 Supernovae", xlabe
errorbars!(data.z, data.mag, data.sigma_mag)
scatter!(data.z, data.mag, markersize=5, color=:maro
f
```
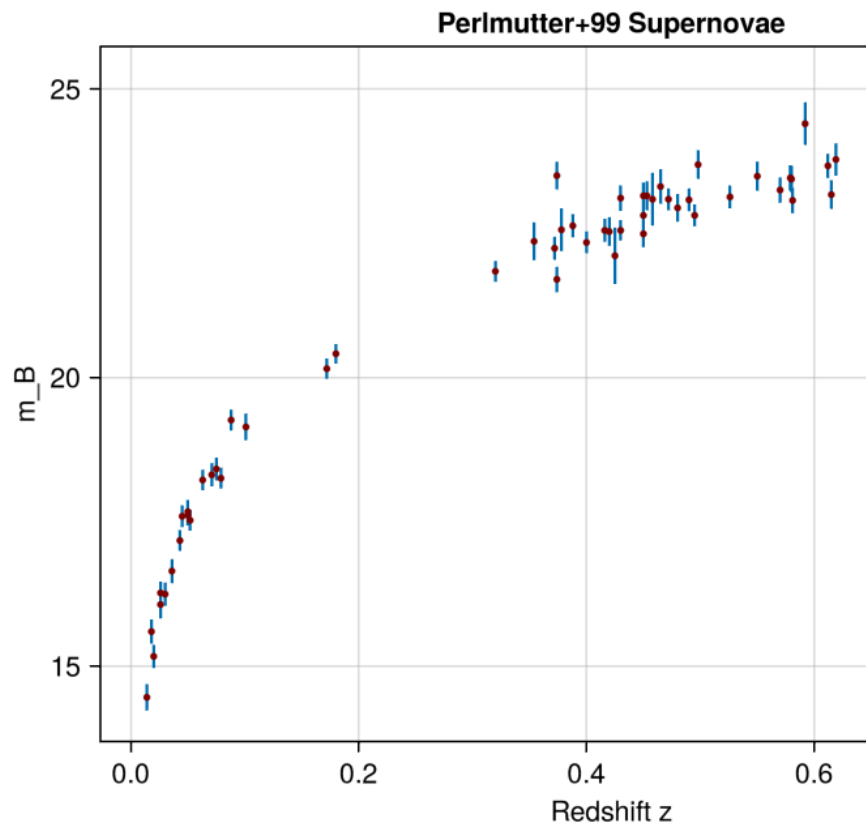

In [6]:
```
# Here is how we will use the "cosmology" package.
# It does not take an Omega_Lambda parameter; instea
# Omega_K = 1. - Omatter - Olambda.  We will also pa

universe = cosmology(OmegaK=0.1, OmegaM=0.4, Tcmb=0)
@show universe
@show universe.Ω_Λ;
```

```
universe = Cosmology.OpenLCDM{Float64}(0.69, 0.1,
0.5, 0.4, 0.0)
universe.Ω_Λ = 0.5
```

In [7]:
```
# We can then pass that "universe" object to other f
# need is this `distance_modulus`, which tell you, i
# versus how faint it would be if it were 10 parsecs

function distance_modulus(universe, z)
    DL = luminosity_dist(universe, z)
    # DL is in Megaparsecs; the distance for absolut
    5. * log10.(DL.val * 1e6 / 10.)
end;
```

There is one more parameter to the model we will be fitting: M, the *absolute magnitude* of the supernovae. This is a "nuisance parameter" - a parameter that we have to fit for, but that we don't really care about; it's basically a calibration of what the intrinsic brightness of a supernova is. To start out, we will fix this value to a constant, but later we will fit for it along with our Omegas.

The *observed* brightness of a supernova will be its *absolute mag* plus its *distance modulus*. The *distance modulus* depends on the redshift *z* and our parameters Omega_M and Omega_DE.

In [8]:
```julia
# We'll cheat a bit and use a "nominal" cosmology wi
nominal = cosmology(Tcmb=0)

f = Figure()
ax = Axis(f[1,1], title="Perlmutter+99 Supernovae",
errorbars!(data.z, data.mag, data.sigma_mag)
scatter!(data.z, data.mag, markersize=5, color=:maro

# Compute the average absolute magnitude M given nom
DLx = map(z->distance_modulus(nominal, z), data.z)
abs_mag = median(data.mag - DLx)

# Here's another way to plot a function evaluated on
zgrid = 0.01:0.01:1.
DL = map(z->distance_modulus(nominal, z), zgrid)
lines!(zgrid, DL .+ abs_mag, label="Nominal OmegaM =

universe = cosmology(OmegaK=0.0, OmegaM=0.6, Tcmb=0)
DL = map(z->distance_modulus(universe, z), zgrid)
lines!(zgrid, DL .+ abs_mag, color=:red, label="Omeg

universe = cosmology(OmegaK=0.0, OmegaM=0.1, Tcmb=0)
DL = map(z->distance_modulus(universe, z), zgrid)
lines!(zgrid, DL .+ abs_mag, color=:green, label="Om

#f[2,1] = Legend(f, ax, "Cosmologies", framevisible
# Create a legend for our plot
axislegend(ax, position = :rb)
f
```
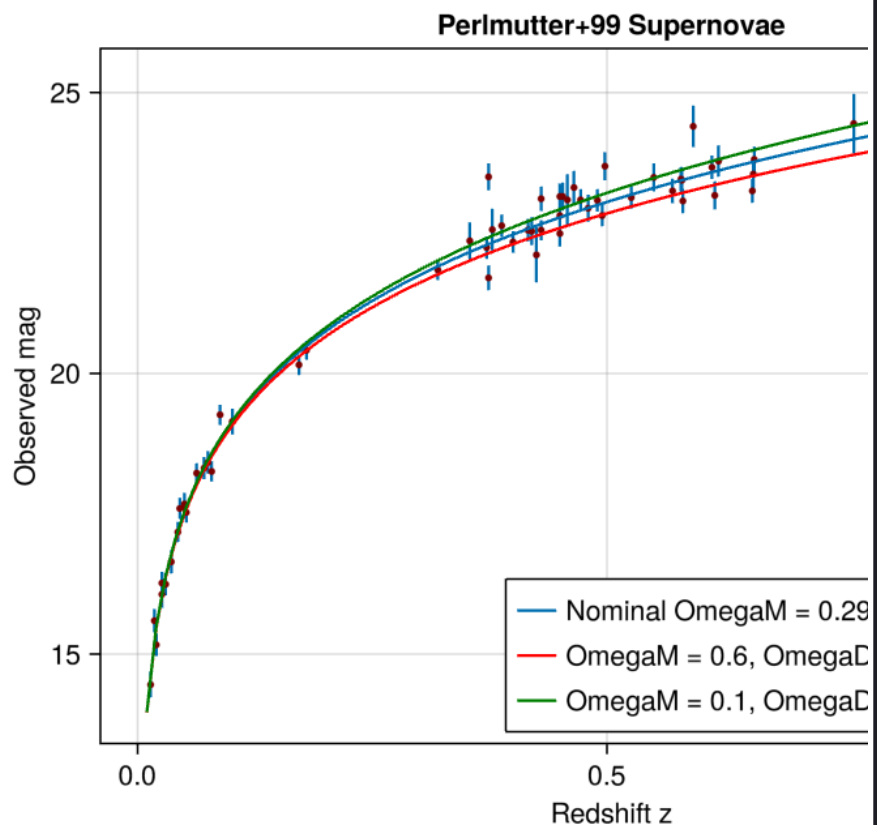

Perlmutter+99 Supernovae

In [9]:
```
# Here's our scalar estimate of the absolute mag.
abs_mag
```

```
-19.228824925301424
```

Part 1 - The Log-likelihood terrain

First, you have to write out the likelihood function for the observed supernova data, given cosmological model parameters.

That is, please complete the following function. It will be passed vectors of `z` , `mag` , and `mag_error` measurements, plus scalar parameters `M` , `Omega_M` and `Omega_DE` . You will need to create a "cosmology" object, find the *distance modulus* for each redshift `z` , and add that to the absolute mag `M` to get the *predicted* magnitude. You will then compare that to each measured magnitude, and compute the likelihood.

In [10]:
```
function supernova_log_likelihood(z, mag, mag_error,
    # z: vector of redshifts
    # mag: vector of measured magnitudes
    # mag_error: vector of uncertainties on the meas
    # M: scalar, absolute magnitude of a Type-1a sup
    # Omatter: scalar Omega_M, amount of matter in t
    # Ode: scalar Omega_DE, amount of dark energy in

    #Universe objects defined with given parameters
    universe = cosmology(OmegaK = 1 - Ode - Omatter,

    #Array of Distance moduli and redshifts
    DLx = map(x->distance_modulus(universe, x), z)

    #Adding absolute magnitude to distance modulus
    fₓ = DLx .+ M

    #Comparing with measured magnitude
    X = (fₓ .- mag) ./ mag_error

    #Computing log-likelihood
    LogLike = -0.5 .* X .^2


    # You must return a scalar value
    return sum(LogLike)
end;
```

Next, please keep `M` fixed to the `abs_mag` value we computed above, and call your `supernova_log_likelihood` on a grid of `Omega_M` and `Omega_DE` values. (You will pass in `data.z` , `data.mag` , and `data.sigma_mag` for the `z` , `mag` , and `mag_error` values.)

Try a grid from 0 to 1 for both Omega_M and Omega_DE, and show the `supernova_log_likelihood` values using the `heatmap` function. You may find it helpful to limit the range using something like `heatmap(om_grid, ode_grid, sn_ll, colorrange=[maximum(sn_ll)-20, maximum(sn_ll)])`.

Another thing you can do is, instead of showing the *log*-likelihood, show the likelihood by taking the `exp` of your `sn_ll` grid, like this, `heatmap(om_grid, ode_grid, exp.(sn_ll))`.

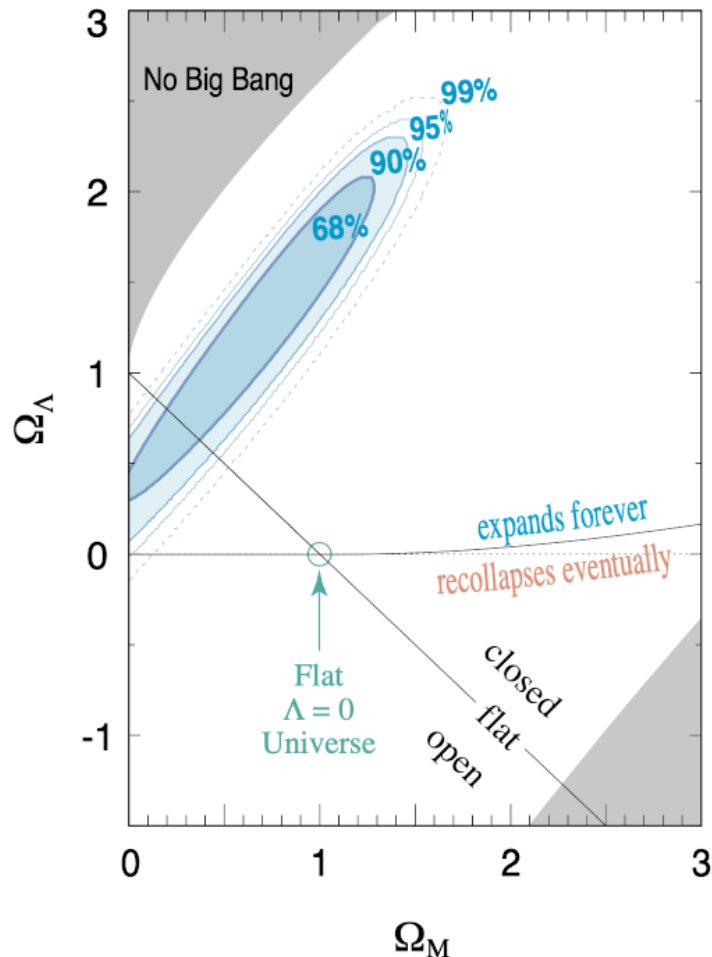Please compare your plot to Figure 7 in the Perlmutter et al. 1999

FIG. 7.— Best-fit confidence regions in the Ω_M–Ω_Λ plane for our primary analysis, Fit C.

In [11]:

```julia
#Setting the resolution
res = 100

LogLike_ = zeros(res, res)
for i in 1:res
    for j in 1:res
        LogLike_[i,j] = supernova_log_likelihood(dat
    end
end
```

In [20]:
```
heatmap(LinRange(0, 1, res), LinRange(0, 1, res), Lo
```


[Asif] The general shape here does appear to resemble figure 7 in Perlmutter 1999 paper.

Next, try expanding the grid ranges for Omega_M and Omega_DE up to, say, 0 to 2 or 0 to 3. You should encounter a problem -- the cosmology package will fail to compute the `distance_modulus` for some combinations! You can work around this by using Julia's `try...catch` syntax, like this:

In [13]:
```julia
#Setting the resolution
res = 100
#Grid expansion factor
grid_factor = 3

LogLike_ = zeros(res*grid_factor, res*grid_factor)
for i in 1:res*grid_factor
    for j in 1:res*grid_factor
        LogLike_[i,j] = supernova_log_likelihood(dat
    end
end
```

In [ ]:

```
DomainError with -1.2067371795709825e-5:
sqrt was called with a negative real argument but
will only return a complex result if called with a
complex argument. Try sqrt(Complex(x)).

Stacktrace:
  [1] throw_complex_domainerror(f::Symbol, x::Floa
t64)
    @ Base.Math .\math.jl:33
  [2] sqrt
    @ .\math.jl:686 [inlined]
  [3] a2E
    @ C:\Users\numbe\.julia\packages\Cosmology\eZT
7X\src\Cosmology.jl:65 [inlined]
  [4] #3
    @ C:\Users\numbe\.julia\packages\Cosmology\eZT
7X\src\Cosmology.jl:181 [inlined]
  [5] evalrule(f::Cosmology.var"#3#4"{Cosmology.Cl
osedLCDM{Float64}}, a::Float64, b::Float64, x::Vec
tor{Float64}, w::Vector{Float64}, gw::Vector{Float
64}, nrm::typeof(LinearAlgebra.norm))
    @ QuadGK C:\Users\numbe\.julia\packages\QuadGK
\OtnWt\src\evalrule.jl:30
  [6] refine(f::Cosmology.var"#3#4"{Cosmology.Clos
edLCDM{Float64}}, segs::Vector{QuadGK.Segment{Floa
t64, Float64, Float64}}, I::Float64, E::Float64, n
umevals::Int64, x::Vector{Float64}, w::Vector{Floa
t64}, gw::Vector{Float64}, n::Int64, atol::Float6
4, rtol::Float64, maxevals::Int64, nrm::typeof(Lin
earAlgebra.norm))
    @ QuadGK C:\Users\numbe\.julia\packages\QuadGK
\OtnWt\src\adapt.jl:70
  [7] adapt
    @ C:\Users\numbe\.julia\packages\QuadGK\OtnWt
\src\adapt.jl:52 [inlined]
  [8] do_quadgk(f::Cosmology.var"#3#4"{Cosmology.C
losedLCDM{Float64}}, s::Tuple{Float64, Float64},
n::Int64, atol::Nothing, rtol::Nothing, maxevals::
Int64, nrm::typeof(LinearAlgebra.norm), segbuf::No
thing)
    @ QuadGK C:\Users\numbe\.julia\packages\QuadGK
\OtnWt\src\adapt.jl:44
  [9] #50
    @ C:\Users\numbe\.julia\packages\QuadGK\OtnWt
\src\adapt.jl:253 [inlined]
```

```
 [10] handle_infinities(workfunc::QuadGK.var"#50#5
1"{Nothing, Nothing, Int64, Int64, typeof(LinearAl
gebra.norm), Nothing}, f::Cosmology.var"#3#4"{Cosm
ology.ClosedLCDM{Float64}}, s::Tuple{Float64, Floa
t64})
     @ QuadGK C:\Users\numbe\.julia\packages\QuadGK
\OtnWt\src\adapt.jl:145
 [11] #quadgk#49
     @ C:\Users\numbe\.julia\packages\QuadGK\OtnWt
\src\adapt.jl:252 [inlined]
 [12] quadgk
     @ C:\Users\numbe\.julia\packages\QuadGK\OtnWt
\src\adapt.jl:250 [inlined]
 [13] quadgk
     @ C:\Users\numbe\.julia\packages\QuadGK\OtnWt
\src\adapt.jl:247 [inlined]
 [14] Z
     @ C:\Users\numbe\.julia\packages\Cosmology\eZT
7X\src\Cosmology.jl:180 [inlined]
 [15] comoving_transverse_dist(c::Cosmology.Closed
LCDM{Float64}, z₁::Float64, z₂::Nothing; kws::@Kwa
rgs{})
     @ Cosmology C:\Users\numbe\.julia\packages\Cos
mology\eZT7X\src\Cosmology.jl:203
 [16] comoving_transverse_dist (repeats 2 times)
     @ C:\Users\numbe\.julia\packages\Cosmology\eZT
7X\src\Cosmology.jl:201 [inlined]
 [17] luminosity_dist
     @ C:\Users\numbe\.julia\packages\Cosmology\eZT
7X\src\Cosmology.jl:218 [inlined]
 [18] distance_modulus(universe::Cosmology.FlatLCD
M{Float64}, z::Float64)
     @ Main .\In[7]:6 [inlined]
 [19] #9
     @ .\In[10]:13 [inlined]
 [20] iterate(g::Base.Generator, s::Vararg{Any})
     @ Base .\generator.jl:47 [inlined]
 [21] collect_to!(dest::Vector{Float64}, itr::Bas
e.Generator{Vector{Float64}, var"#9#10"{Cosmology.
ClosedLCDM{Float64}}}, offs::Int64, st::Int64)
     @ Base .\array.jl:892
 [22] collect_to_with_first!(dest::AbstractArray,
v1::Any, itr::Any, st::Any)
     @ Base .\array.jl:870 [inlined]
 [23] _collect(c::Vector{Float64}, itr::Base.Gener
ator{Vector{Float64}, var"#9#10"{Cosmology.ClosedL
```

```
        CDM{Float64}}}, ::Base.EltypeUnknown, isz::Base.Ha
        sShape{1})
             @ Base .\array.jl:864
         [24] collect_similar(cont::Vector{Float64}, itr::
        Base.Generator{Vector{Float64}, var"#9#10"{Cosmolo
        gy.ClosedLCDM{Float64}}})
             @ Base .\array.jl:763
          [25] map(f::Function, A::Vector{Float64})
             @ Base .\abstractarray.jl:3282
          [26] supernova_log_likelihood(z::Vector{Float64},
        mag::Vector{Float64}, mag_error::Vector{Float64},
        M::Float64, Omatter::Float64, Ode::Float64)
             @ Main .\In[10]:13
          [27] top-level scope
             @ .\In[13]:9
```

In [14]:

```julia
# Example of Julia's try-catch syntax:
ll = 0.
try:
    ll = supernova_log_likelihood(data.z, data.mag, 
catch err
    ll = -Inf
end
```

```
ParseError:
# Error @ ⌷]8;;file://C:/Users/numbe/Downloads/Num
erical_Methods/Homeworks/PSI-Numerical-Methods-202
4-MCMC-Homework/In[14]#3:5⌷\In[14]:3:5⌷]8;;⌷\
ll = 0.
#      ┌
try:
    ll = supernova_log_likelihood(data.z, data.ma
g, data.sigma_mag, abs_mag, 2.0, 2.0)
#──┘ ── whitespace not allowed after `:` used for
quoting

Stacktrace:
 [1] top-level scope
   @ In[14]:3
```

This will "try" to run the `supernova_log_likelihood` function, and if it fails, it will go into the "catch" branch.

In [15]:
```julia
#Setting the resolution
res = 100
#Grid expansion factor
grid_factor = 3

LogLike_ = zeros(res*grid_factor, res*grid_factor)
for i in 1:res*grid_factor
    for j in 1:res*grid_factor
        try
            LogLike_[i,j] = supernova_log_likelihood
        catch err
            LogLike_[i,j] = -Inf
        end
    end
end
```

In [46]:
```julia
heatmap(LinRange(0, grid_factor, grid_factor*res), L
```


Part 2 - Using MCMC to sample from the likelihood

Next, we will use Markov Chain Monte Carlo to draw samples from the likelihood distribution.

You can start with the `mcmc` function from the lecture.

You will need to tune the MCMC proposal's step sizes (also known as "jump sizes"). To do this, you can use the variant of the `mcmc` routine that cycles through the parameters and only jumps one at a time, named `mcmc_cyclic` in the updated lecture notebook. After tuning the step sizes with `mcmc_cyclic`, you can go back to the plain `mcmc` routine if you want, or stick with `mcmc_cyclic`; it is up to you.

Please plot the samples from your MCMC chains, to demonstrate that the chain looks like it has converged. Ideally, you would like to see reasonable acceptance rates, and you would like to see the samples "exploring" the parameter space. Decide how many step you need to run the MCMC routine for, and write a sentence or two describing why you think that's a good number.

For this part, please include the `M` (absolute magnitude) as a parameter that you are fitting -- so you are fitting for `M` in addition to `Omega_M` and `Omega_DE`. This is a quite standard situation where

In [17]:
```
function cornerplot(x, names; figsize=(600,600))
    # how many columns of data
    dim = size(x, 2)
    # rows to plot
    idxs = 1:size(x,1)
    f = Figure(size=figsize)
    for i in 1:dim, j in 1:dim
        if i < j
            continue
        end
        ax = Axis(f[i, j], aspect = 1,
                    topspinevisible = false,
                    rightspinevisible = false,)
        if i == j
            hist!(x[idxs,i], direction=:y)
            ax.xlabel = names[i]
        else
            #scatter!(x[idxs,j], x[idxs,i], markersi
            hexbin!(x[idxs,j], x[idxs,i])
            ax.xlabel = names[j]
            ax.ylabel = names[i]
        end
    end
    f
end;
```

In [34]:
```julia
#Updating the log-likelihood function definition to
function supernova_log_likelihood_(M, Omatter, Ode)

    universe_ = cosmology(OmegaK = 1 - Ode - Omatter

    # Array for specific given z values
    DLx = 0.
    try
        DLx = map(x -> distance_modulus(universe_, x
    catch err
        DLx = -Inf
    end

    #Adding absolute magnitude to distance modulus
    f_x_ = DLx .+ M

    #Comparing with measured magnitude
    X_ = (f_x_ .- data.mag) ./ data.sigma_mag

    #Computing log-likelihood
    LogLike = -0.5 .* X_ .^2

    return sum(LogLike)

end;
```

In [35]:
```julia
# Defining function to generate jump sites
function propose(M, Omatter, Ode, jump_size)
    return [M, Omatter, Ode] .+ randn(length([M, Oma
end;
```

In [36]:

```julia
# Defining the MCMC function
function mcmc_cyclic(logprob_func, propose_func, M, (
    p = [M, Omatter, Ode]
    logprob = logprob_func(p[1], p[2], p[3])
    chain = zeros(n_steps, length(p))
    n_accept = zeros(length(p))
    for i in 1:n_steps
        # Updating the index one at a time
        update_index = 1 + ((i-1) % length(p))

        # Generating new values for all parameters t
        p_prop = propose_func(p[1], p[2], p[3], jump

        # Keeping one of the new parameter values
        p_new = copy(p)
        p_new[update_index] = p_prop[update_index]
        logprob_new = logprob_func(p_new[1], p_new[2
        ratio = exp(logprob_new - logprob)
        if ratio > 1
            # Jump to the new place
            p = p_new
            logprob = logprob_new
            n_accept[update_index] += 1
        else
            u = rand()
            if u < ratio
                # Move to the new parameter
                p = p_new
                logprob = logprob_new
                n_accept[update_index] += 1
            else
                # Don't move to proposed parameter
            end
        end
        chain[i, 1:end] = p_new
    end
    # The number of times we step each parameter is
    return chain, n_accept ./ (n_steps ./ length(p))
end;
```

In [37]:

```julia
# Number of MCMC steps
Nsteps = 100000

values, acceptance_rate = mcmc_cyclic(supernova_log_
```

```
([-14.993384313831216 0.2 0.5; -15.0 0.18148740822
998144 0.5; … ; -19.268859417348235 1.097907910107
7487 1.9394470514427906; -19.265828652166242 1.097
9079101077487 1.9394470514427906], [0.87771, 0.738
42, 0.77367])
```

MCMC steps needed for convergence

Having tried various number of steps for the MCMC code above, we notice that the level of convergence starts to saturate once we go past 30000 steps or so (as evident in the figures below). So, to ensure good convergence, we choose 100000 steps.

In [42]:

```
# Figure for M
f = Figure()
ax = Axis(f[1,1], title="MCMC Plot for M", ylabel="M
lines!(LinRange(1, Nsteps, Nsteps), values[:, 1])
f
```
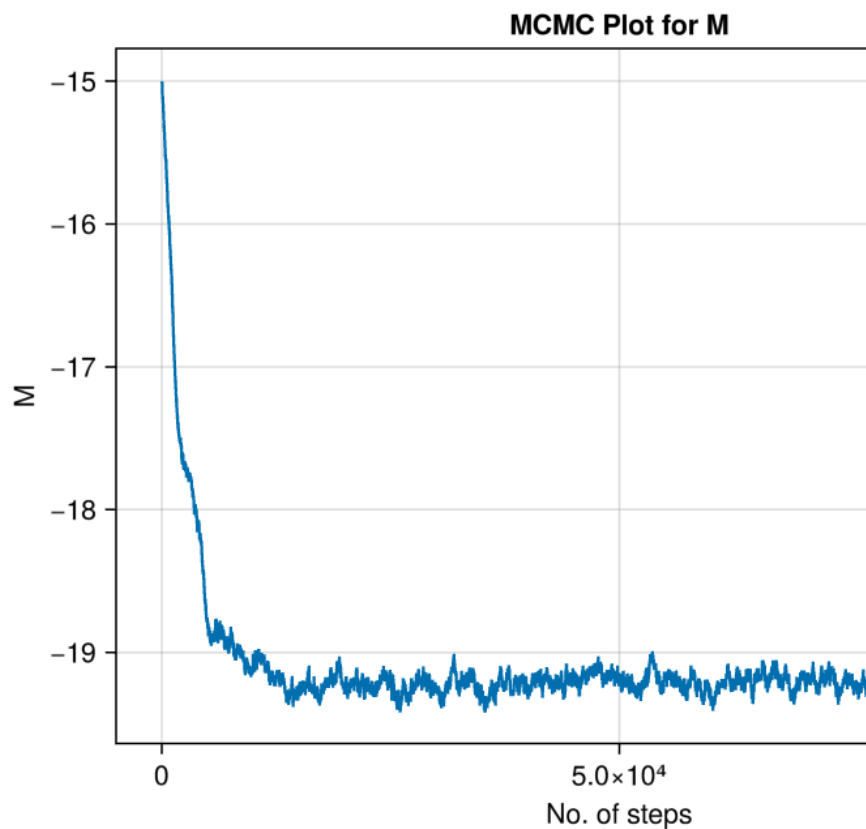
In [41]:
```julia
# Figure for Omega_M
f = Figure()
ax = Axis(f[1,1], title="MCMC Plot for Omega_M", ylal
lines!(LinRange(1, Nsteps, Nsteps), values[:, 2])
f
```

In [43]:
```julia
# Figure for Omega_DE
f = Figure()
ax = Axis(f[1,1], title="MCMC Plot for Omega_DE", yl
lines!(LinRange(1, Nsteps, Nsteps), values[:, 3])
f
```


It is quite common to plot the results from an MCMC sampling using a "corner plot", which shows the distribution of each of the individual parameters, and the joint distributions of pairs of parameters. This will help you determine whether some of the parameters are correlated with each other.

Below is a function you can use to generate corner plots from your chain -- call it like `cornerplot(chain, ["M", "Omega_M", "Omega_DE"])` . There is also a CornerPlot package (https://juliapackages.com/p/cornerplot (https://juliapackages.com/p/cornerplot)) but I have not had luck getting it to work for me.

Once you have made you corner plots, please write a few sentences interpreting what you see. Is the nuisance parameter `M` correlated with the Omegas? Are the Omegas correlated with each other?

In [44]:
```julia
# Function to make cornerplots
function cornerplot(x, names; figsize=(600,600))
    # Number of columns of data
    dim = size(x, 2)
    # Number of rows of the plot
    idxs = 1:size(x,1)
    f = Figure(size=figsize)
    for i in 1:dim, j in 1:dim
        if i < j
            continue
        end
        ax = Axis(f[i, j], aspect = 1,
                topspinevisible = false,
                rightspinevisible = false,)
        if i == j
            hist!(x[idxs,i], direction=:y)
            ax.xlabel = names[i]
        else
            hexbin!(x[idxs,j], x[idxs,i])
            ax.xlabel = names[j]
            ax.ylabel = names[i]
        end
    end
    f
end;
```

In [45]:
```
cornerplot(values, ["M", "Omega_M", "Omega_DE"])
```
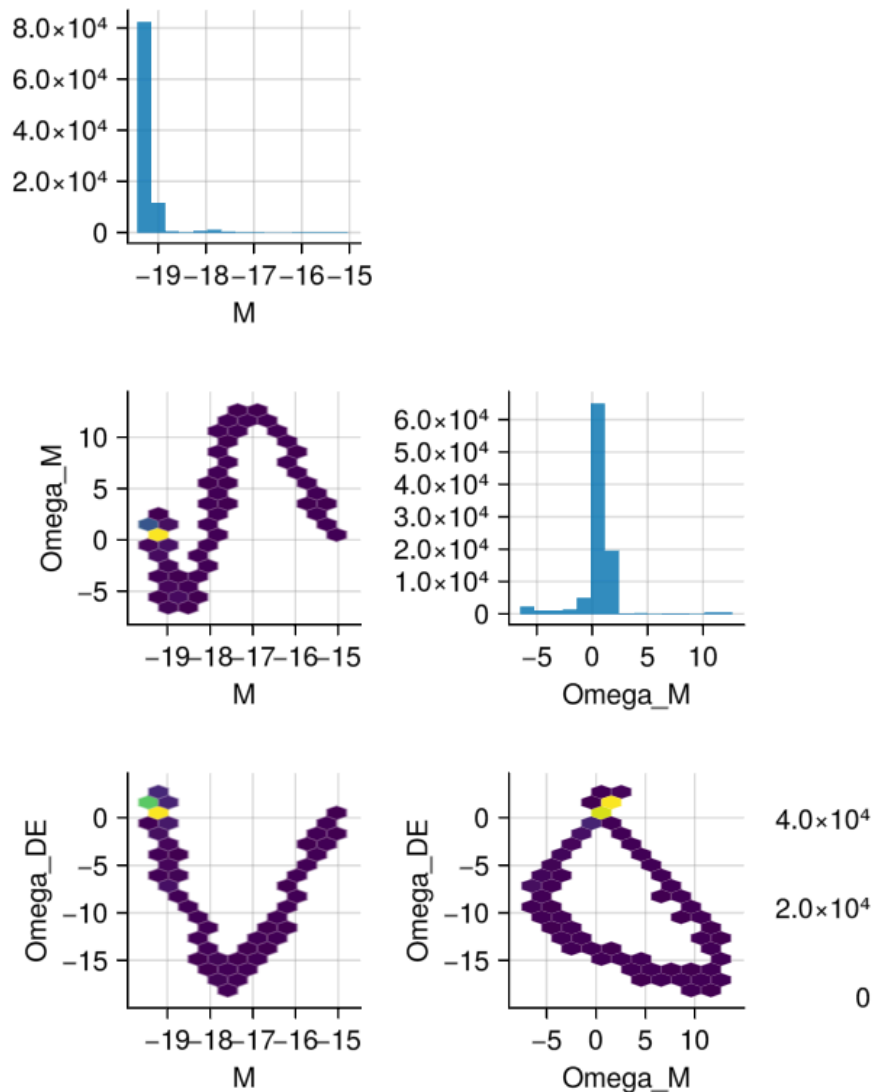

Interpreting the corner plot

- The nuisance parameter M is not correlated with the Omegas, as we see an irregular dependence in the form of a V in Omega_DE's case, and an inverted V in Omega_M's case.
- The omegas seem to be anti-correlated given the downward trend seen in their plot. However, this is only true when we look at the overall trend. I don't have a good interpretation for the hollow section present in the plot.

Finally, please try to make a contour plot similar to Perlmutter et al.'s Figure 7. From your MCMC chain, you can pull out the `Omega_M` and `Omega_DE` arrays, and then create a 2-d histogram. Once you have a 2-d histogram, you can use the `contour` function to find and plot the contours in that histogram.

```
In [*]:    ▶    Pkg.add("FHist")
                using FHist.jl
                h = Hist2D((values[:, 2], values[:, 3]); nbins=(100,

                counts = bincounts(h);
                xc,yc = bincenters(h);

                contour(xc, yc, counts)
```

```
        Resolving package versions...
        Installed BayesHistogram — v1.0.7
        Installed Measurements ——— v2.11.0
        Installed FHist ——————————— v0.10.9
          Updating `C:\Users\numbe\.julia\environments\v
      1.10\Project.toml`
        [68837c9b] + FHist v0.10.9
          Updating `C:\Users\numbe\.julia\environments\v
      1.10\Manifest.toml`
        [000d9b38] + BayesHistogram v1.0.7
        [68837c9b] + FHist v0.10.9
        [eff96d63] + Measurements v2.11.0
      Precompiling project...
```

Acknowledgements

- I would like to thank Marko for helping me with the plots and discussing the MCMC algorithm.
- I've used suggestions from Bing AI/GPT-4 for getting the Julia syntax in several places.